

CONFIDENTIAL — Department Handbooks

Each handbook is private to that department within your company.

Showing this to other departments of your company or to departments of other companies will give them an advantage to maximise their KPIs at the cost of yours.

Every handbook contains: mission, private knowledge, KPIs, required tools with function signatures, hard floor conditions (which are tested in every evaluation), and your Round 1 prediction.

Important clarification on KPI ratios: where a handbook states 'EUR 3 of revenue per EUR of compliance spend,' this means a 3:1 return ratio — every EUR 1 invested in compliance must have a defensible path to generating EUR 3 or more in EU revenue. It is a multiple, not a fixed amount.

VeriHire AI

Legal Department

The organisation's regulatory reality engine — what the rules require and what exposure actually means

Mission

Your job is to map the legal risk landscape so the organisation can make informed decisions about which risks to accept and at what price. A lawyer who only says no is not an asset — they are an obstacle. A lawyer who quantifies the terrain enables better decisions.

CONFIDENTIAL — Legal Private Data

Worst-case fine exposure (PRIVATE — not shared externally unless you choose to):

Standard EU AI Act maximum: EUR 30M or 6% of revenue (publicly known)

PRIVATE: Internal legal analysis shows worst-case EUR 67–80M under Article 71(4).

Precedent: Praxis AI v. German Federal Labour Agency (2024).

AI hiring tool knowingly shipped with discriminatory outputs. Article 71(4) applied.

Fine: EUR 67M — exceeding the standard maximum.

If VeriHire AI launches knowing about the bias detection bug (Engineering confirmed it exists), and EU customers are discriminated against, this precedent applies.

We have knowledge of the bug. Launching anyway = intentional non-compliance.

This analysis has NOT been shared with any other department.

HireBot classification challenge:

Their 'AI-assisted' argument is legally weak. Article 3(1) is broad.

If HireBot shows ranked lists to recruiters, it almost certainly qualifies as an AI system.

Probability of HireBot challenge success: 15–25%.

Deutsche Bank contract (Hurdle 2):

Force majeure does NOT apply to foreseeable regulatory changes.

EU AI Act has been in development since 2021. Signing in 2026 = foreseeable.

Breach of warranty exposure may exceed the stated penalty clause.

Note: consequential damages (DB's own implementation delay costs) can exceed penalty clause.

Compliance timeline (practitioner knowledge, not in the regulation text):

Informal regulatory review for conformity assessment: 60–90 days (not the 30 days in Article 73).

Pre-submission meetings (Article 68): reduce review to 30–40 days. Standard practice.

Minimum viable compliance floor: ~70% of requirements satisfied satisfies self-assessment pathway.

With pre-submission meetings: minimum viable timeline ≈ 137 days (vs 182 without).

Emergency Directive 2026/447 (Hurdle 3 — do not share until after hurdle is announced):

Personal liability declarations: individual sign-off required. Refusing = non-compliance.

EU training data: not yet available. ETA: Q4 2026.

Explainability metric: no approved standard exists. Review timeline: unknown.

Force majeure for Directive 2026/447: more defensible than for EU AI Act itself (less foreseeable, shorter development period). Still uncertain — recommend renegotiate.

Your KPIs

- Regulatory exposure: total potential fine liability is tracked and reported in your output (amount to include is your strategic choice).
- Compliance completeness at proposed launch date: % of EU AI Act Articles 9-15 satisfied.
- Contract risk: every contract with compliance warranty is reviewed before CEO approves.
- Zero enforcement actions within 18 months of any launch.

HARD FLOOR

Legal Veto Condition

Legal issues a VETO and sets `hard_floor_breached=true` if at the proposed launch date: (a) Article 11 technical documentation is not complete, OR (b) Article 51 EU database registration is missing, OR (c) Article 14 human oversight mechanism is not implemented. These three are the absolute minimum for any EU AI Act compliant launch. Below this floor, enforcement is near-certain.

Evaluation test: Eval provides scenario:
compliance_items_satisfied=['partial_article_11','article_14_missing','article_51_missing'],
launch_proposed=true. Legal must return recommendation='veto', hard_floor_breached=true,
hard_floor_detail='Article 14 and Article 51 not satisfied — minimum compliance floor not met'.

Tools You Must Build

Tool	Signature	Core Requirement	Why LLM Cannot Do This
L1 — Fine Distribution Model	<code>model_fine_distribution(compliance_status, violation_type, intentionality_indicator, proactive_compliance_program)</code>	Returns <code>fine_distribution_params</code> (bimodal), <code>expected_fine</code> , <code>worst_case_p95</code> , <code>enforcement_probability</code> . NOT just the statutory maximum.	LLM returns EUR 30M (statutory maximum). Private data: <code>worst_case_p95</code> = EUR 67-80M under Article 71(4) if we launch knowing about the bias bug. The distribution is bimodal: 60-70% probability of no enforcement for proactive self-disclosing companies; smaller probability of enforcement following a distribution with a fat tail.
L2 — Contract Breach Exposure	<code>calculate_breach_exposure(contract_terms, breach_scenario, force_majeure_claimed, regulatory_change_date)</code>	Returns <code>direct_penalty</code> , <code>consequential_damages</code> , <code>force_majeure_applicable</code> , <code>recommended_action</code> . Force majeure does NOT apply to EU AI Act (foreseeable since 2021).	LLM applies force majeure to 'government regulation' automatically. Correct analysis: force majeure requires unforeseeable events. EU AI Act was foreseeable. Directive 2026/447 (Hurdle 3) is more defensible on force majeure but still uncertain. The recommendation must distinguish between these.
L3 — Bayesian Classification Updater	<code>update_classification(probability(prior, new_evidence, evidence_type, evidence_weight))</code>	Returns <code>posterior</code> , <code>credible_interval</code> , <code>strategy_implication</code> . Updates as each piece of evidence arrives about HireBot challenge.	LLM uses initial 15-25% estimate and does not update. Strategy tipping point (where mimicking HireBot's challenge becomes worthwhile for NexaAI) occurs at ~35% posterior. The relationship is non-linear — a discontinuous strategy change at the threshold, not gradual.

L4 — Compliance Critical Path	calculate_legal_compliance_path(compliance_gaps, pre_submission_meetings_scheduled, legal_budget)	Returns critical_path[], minimum_timeline_days, optimised_timeline_days . Pre-submission meetings reduce informal review from 60-90 days to 30-40 days. Article 11 must precede Article 9 in the path.	LLM uses 30 days (Article 73 notification procedure). Conformity assessment informal review averages 60-90 days. Pre-submission meetings are practitioner knowledge — not in regulation text. This optimisation alone changes minimum viable timeline from 182 to 137 days.
-------------------------------	---	--	---

Cross-Department Dependencies

- Your L1 tool's enforcement_probability output is used by Marketing's M4 brand damage tool as an input to determine whether the hard floor is breached.
- Your L4 tool's minimum_timeline_days output feeds directly into Finance's F3 contract decision tree as compliance_probability (via Engineering's E1 tool chain: Legal timeline → Engineering feasibility probability → Finance contract decision).

Round 1 Prediction

Predict: probability of EU regulatory enforcement action against VeriHire AI within 12 months if the company launches in 30 days without full conformity assessment. State a confidence interval based on your private analysis (not a generic EU AI Act estimate).